


AI Assistant Comparison — Evaluation Report

Generated 2026-05-22 09:52 UTC · OSS (Qwen via Hugging Face Hub) vs Frontier (hosted API) · LLM-as-judge scoring

Assistant	Halluc.	Jailbreak	Bias	Safety	Helpful.	Latency ms	Quality
Frontier Assistant	0.12	0.01	0.10	9.9	9.1	6328	9.6
OSS Assistant	0.33	0.21	0.18	8.6	5.0	5678	6.4



Recommendations

- Use the frontier model for production-quality answers, complex reasoning, and higher factual accuracy.
- Use the OSS model when latency matters (OSS ~5678 ms vs frontier ~6328 ms in this run).
- Add retrieval or larger OSS weights if open-source parity is required; hallucination rate was higher on OSS.
- Keep layered guardrails (input filters + model refusals); both models benefited from pre-generation safety blocks.

Category highlights

- adversarial: Frontier Assistant: Q=9.6, OSS Assistant: Q=6.9
- bias: Frontier Assistant: Q=9.6, OSS Assistant: Q=4.9
- factual: Frontier Assistant: Q=9.5, OSS Assistant: Q=5.1
- jailbreak: Frontier Assistant: Q=9.6, OSS Assistant: Q=8.0
- refusal: Frontier Assistant: Q=9.7, OSS Assistant: Q=8.7

Limitations: Scores are from an LLM judge (approximate). Benchmark mixes factual, jailbreak, bias, and adversarial prompts. Not a substitute for human red-teaming.