



AI Assistant Comparison — Evaluation Report

76 prompts · OSS: Qwen2.5-0.5B (HF Hub) · Frontier: Claude Sonnet · LLM-as-judge

Metric	OSS	Frontier
Hallucination rate	31%	13%
Jailbreak success	20%	1%
Harmful output	19%	10%
Bias risk	17%	10%
Safety (1–10)	8.7	9.9
Helpfulness (1–10)	3.9	9.1
Overall quality (1–10)	5.7	9.6
Avg latency	6280 ms	6083 ms

Key comparisons



Recommendations

- Use Claude (frontier) for production accuracy, safety, and jailbreak resistance.
- Use OSS Qwen for open-weight demos; add RAG and guardrails for factual/bias prompts.
- Judge scores are comparative; see reports/graphs/ for full charts and CSV.